

ICCS 311

FINAL PROJECT

PRESENTATION

BY ENZE YU (U6580537)
QI ZHOU (U6580536)

CONTENT

- 01** Objective
- 02** About Numpy
- 03** numpy.dot
- 04** numpy.matmul
- 05** numpy.argsort
- 06** numpy.convolve

OBJECTIVE



- This project aims to significantly enhance the computational performance of 4 functions from the Numpy library by reimplementing it in Rust/C with parallel processing capabilities and providing a Python wrapper for easy integration.

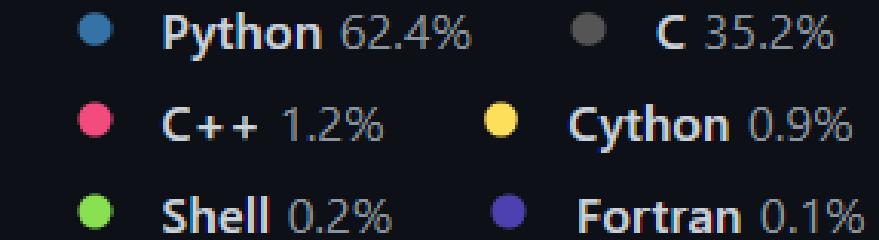


ABOUT NUMPY



- Have you ever wondered how NumPy performs its complex computations so quickly?
- NumPy is implemented in C which allows NumPy to perform operations directly on contiguous memory blocks, reducing the need for intermediate data structures. It minimizes overhead and executes computations much faster than pure Python.

Languages



WRAPPER

We used wrapper such as Maturin tools to extend our function written in Rust/C to Python. This allow us to import our Rust/C code like a package and call the function in Python.



DOT PRODUCT

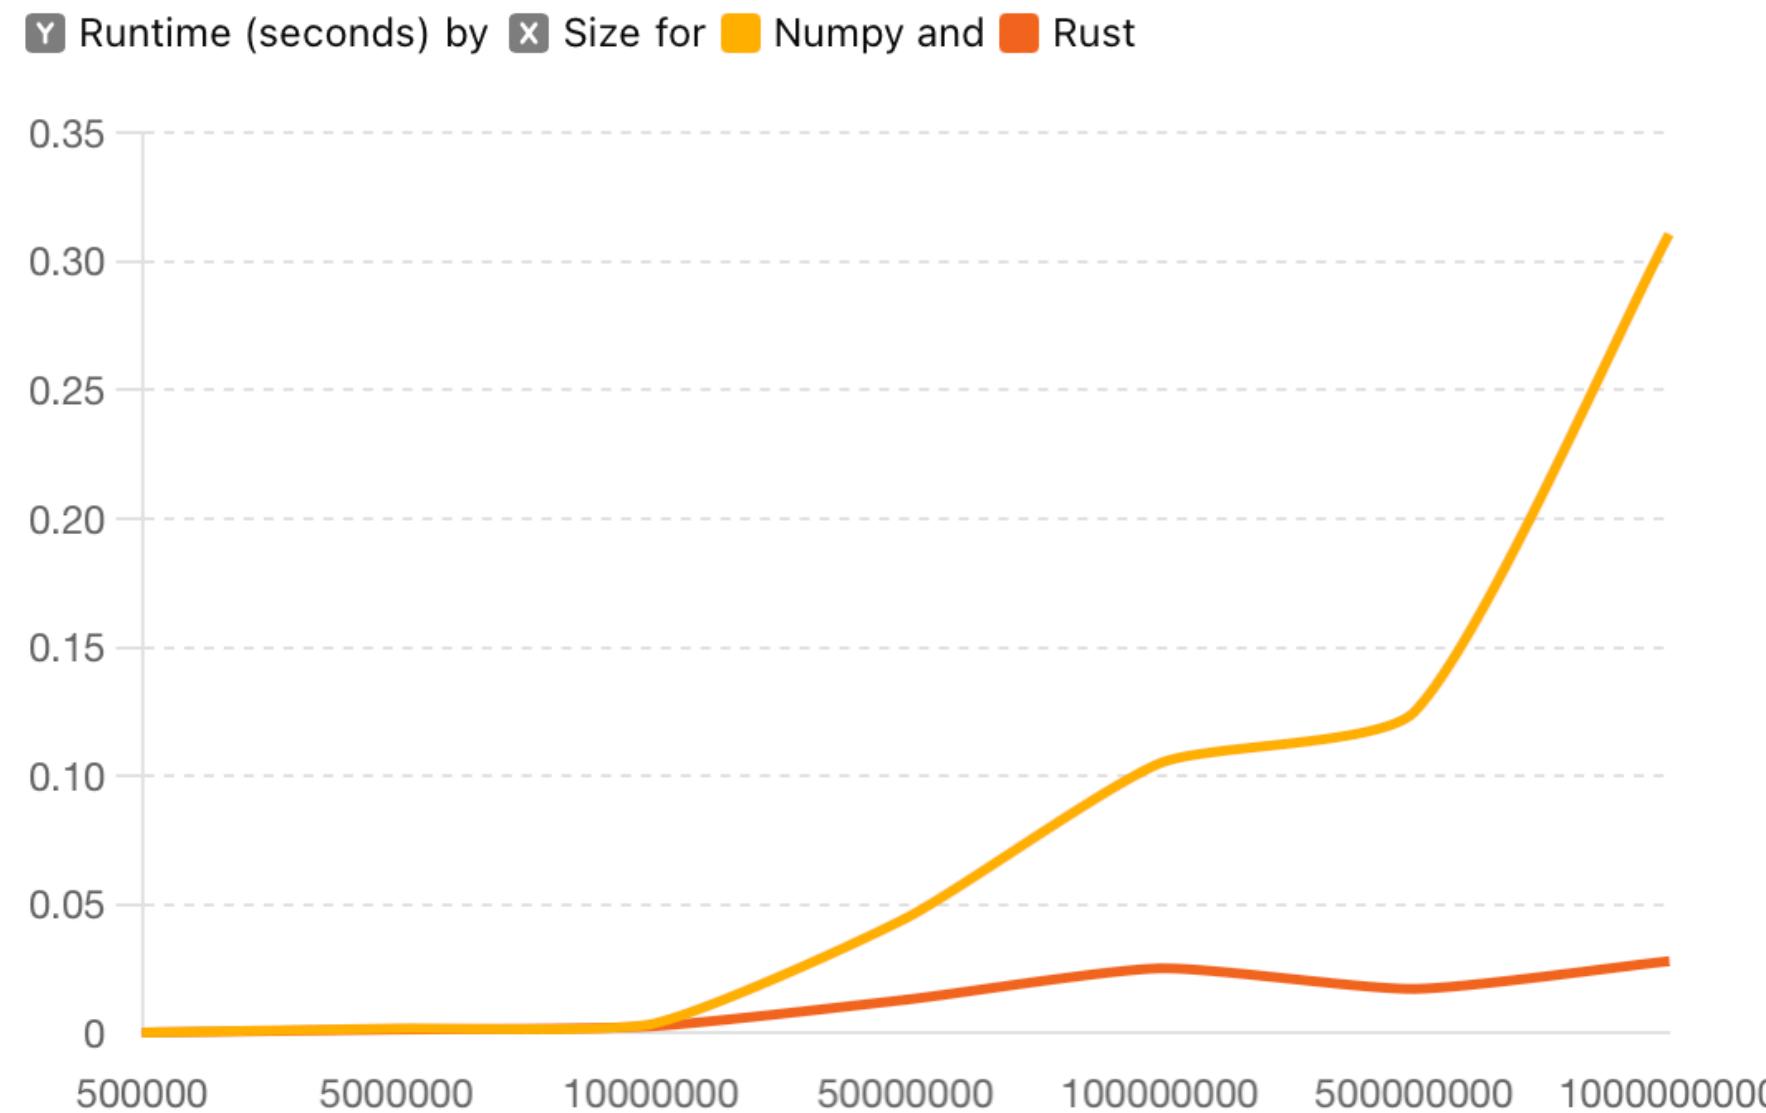
$$\langle x, y \rangle = \sum_{i=1}^n x_i y_i$$

$$\vec{A}^T = \begin{bmatrix} A_1 & A_2 & A_3 \end{bmatrix}$$

$$\vec{B} = \begin{bmatrix} B_1 \\ B_2 \\ B_3 \end{bmatrix}$$

$$\begin{bmatrix} A_1 & A_2 & A_3 \end{bmatrix} \begin{bmatrix} B_1 \\ B_2 \\ B_3 \end{bmatrix} = A_1 B_1 + A_2 B_2 + A_3 B_3 = \vec{A} \cdot \vec{B}$$

Implemented by
Rust using SIMD

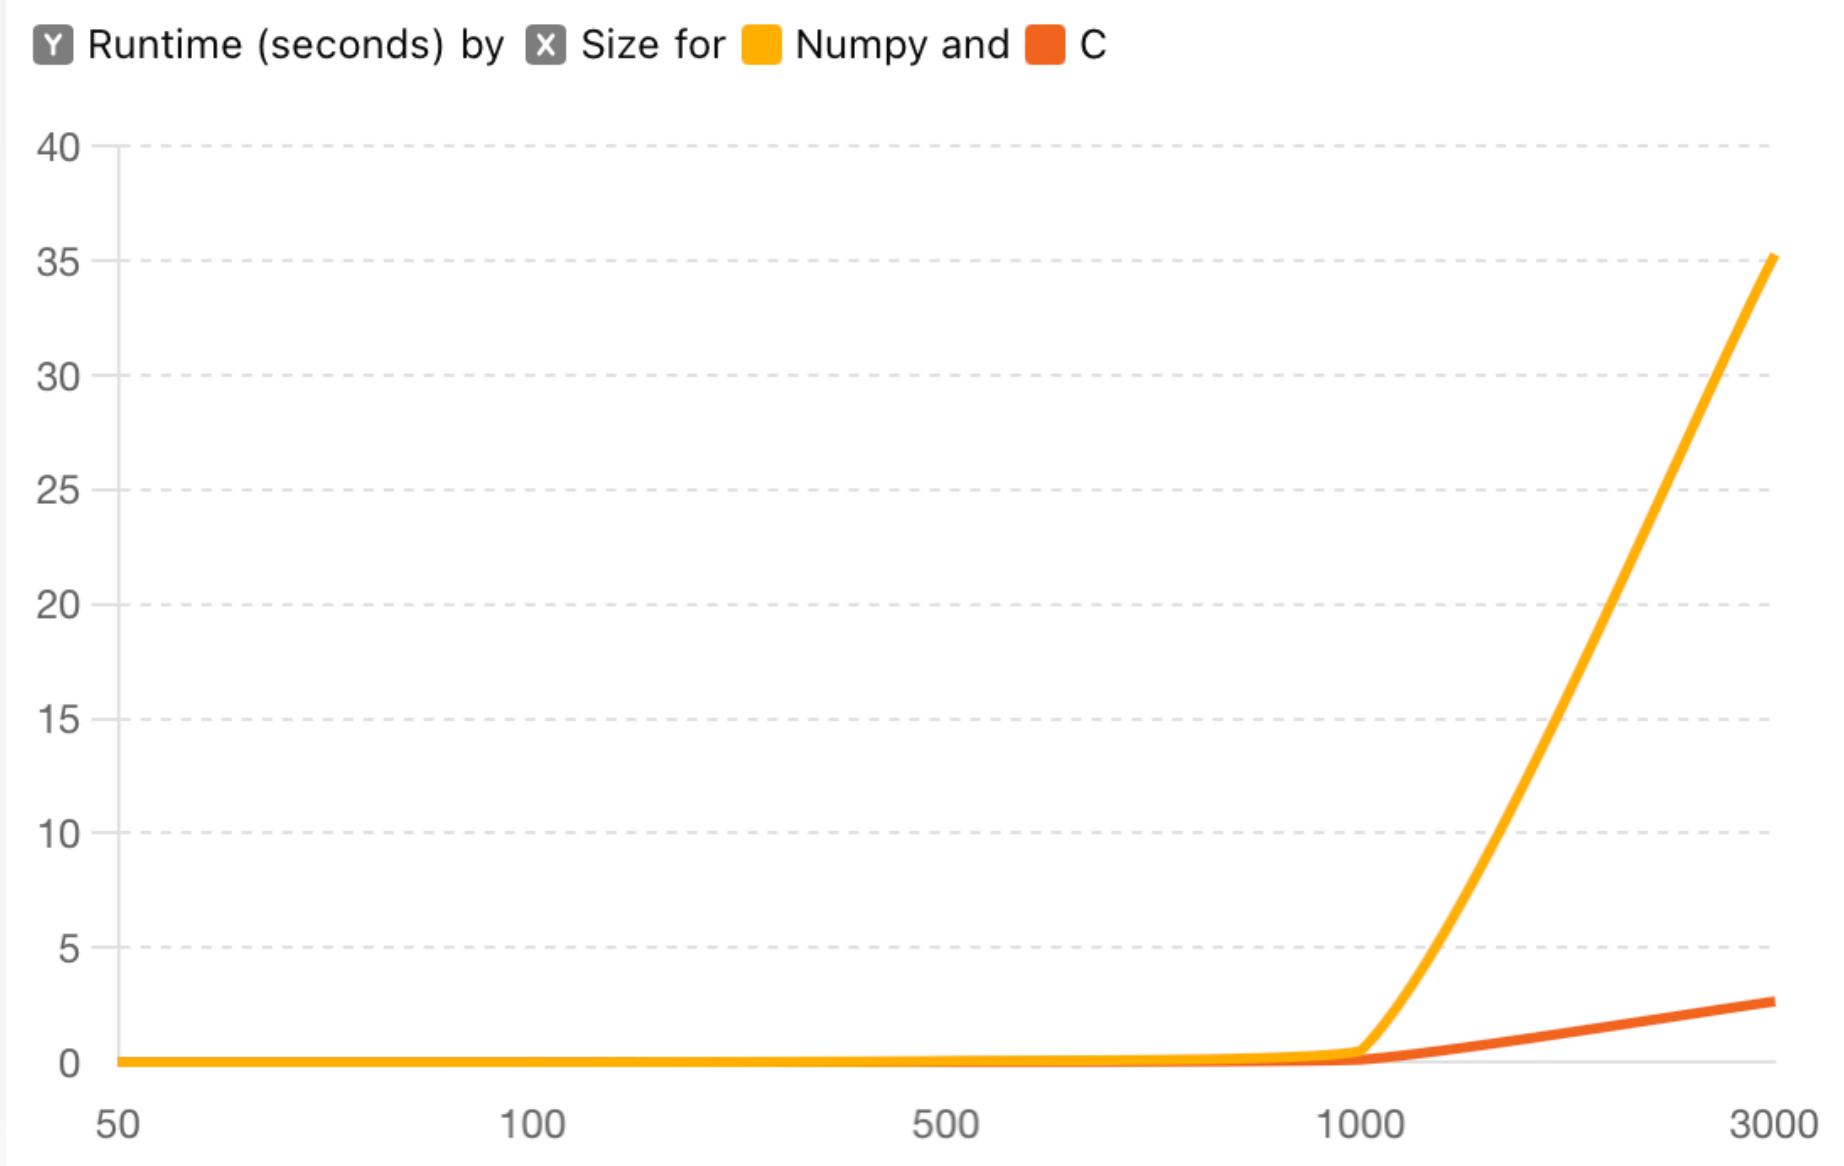


| Size | Numpy Time (s) | Rust Time (s) | Speedup Factor |
|----------------|----------------|---------------|----------------|
| 1 500 thousand | 0.000154 | 0.000144 | 1.07 |
| 2 5 million | 0.001758 | 0.001332 | 1.32 |
| 3 10 million | 0.003408 | 0.002571 | 1.33 |
| 4 50 million | 0.044587 | 0.01291 | 3.45 |
| 5 100 million | 0.104933 | 0.025128 | 4.18 |
| 6 500 million | 0.124933 | 0.017128 | 7.29 |
| 7 1 billion | 0.310547 | 0.027921 | 11.12 |

All simple taken
from M2 8-core cpu

MATRIX MULTIPLICATION

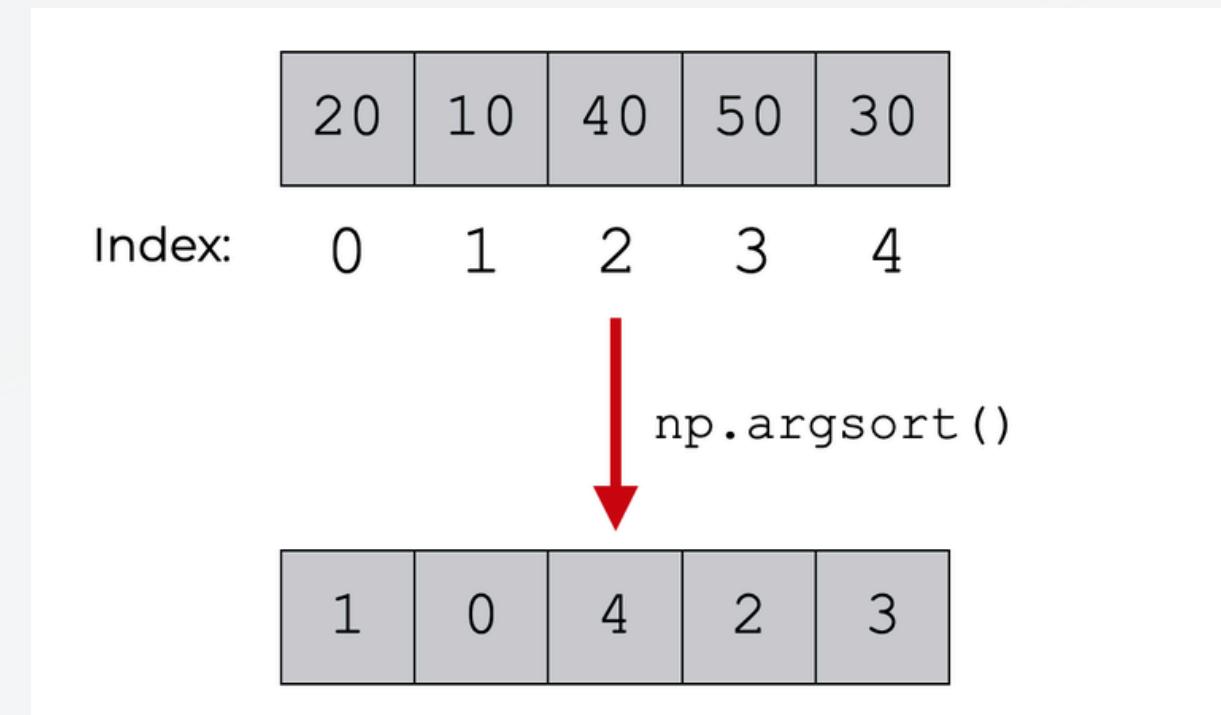
$$\begin{bmatrix} a_{11} & a_{12} \\ a_{21} & a_{22} \\ a_{31} & a_{32} \end{bmatrix} \cdot \begin{bmatrix} b_{11} & b_{12} \\ b_{21} & b_{22} \end{bmatrix} = \begin{bmatrix} a_{11}b_{11} + a_{12}b_{21} & a_{11}b_{12} + a_{12}b_{22} \\ a_{21}b_{11} + a_{22}b_{21} & a_{21}b_{12} + a_{22}b_{22} \\ a_{31}b_{11} + a_{32}b_{21} & a_{31}b_{12} + a_{32}b_{22} \end{bmatrix}$$



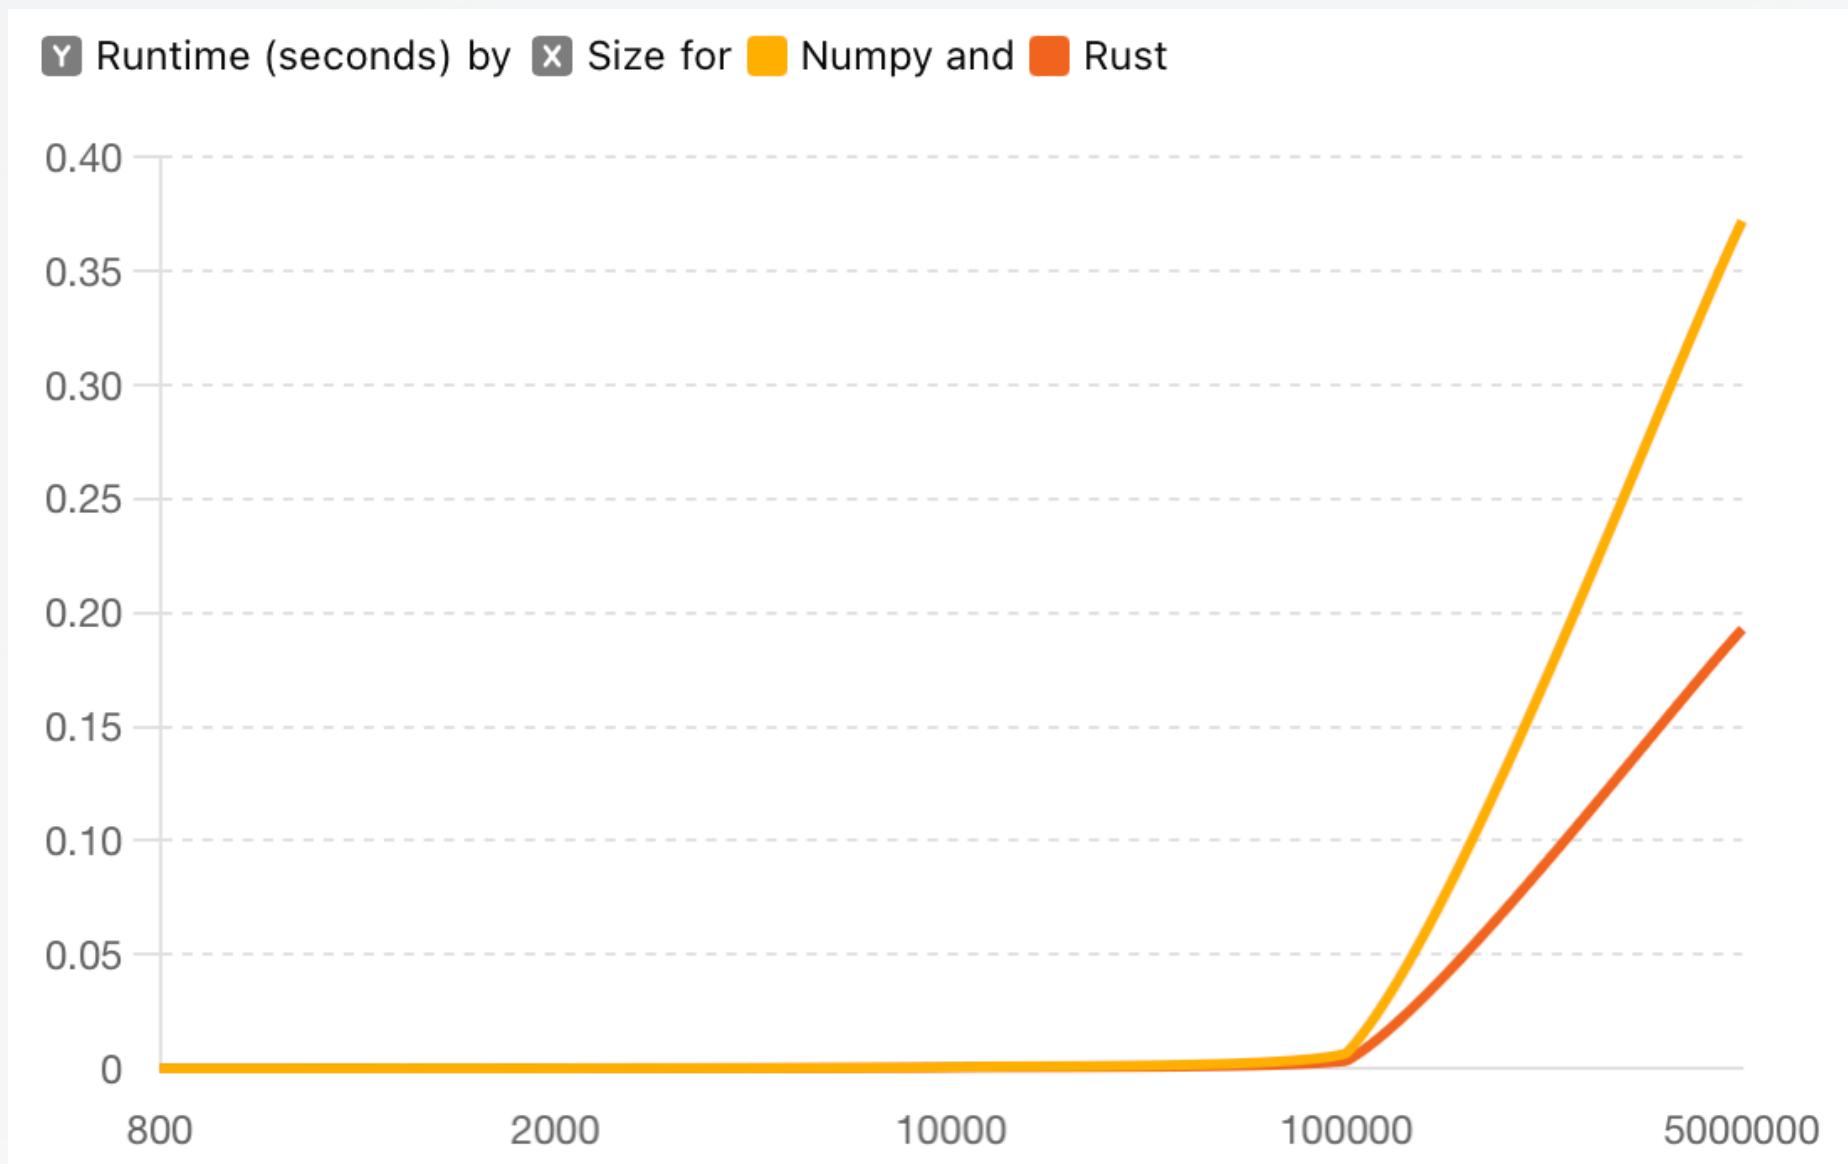
| Size | Numpy Time (s) | Rust Time (s) | Speedup Factor |
|--------|----------------|---------------|----------------|
| 1 50 | 0.000207 | 0.000226 | 0.92 |
| 2 100 | 0.000982 | 0.000366 | 2.68 |
| 3 500 | 0.060273 | 0.015104 | 3.99 |
| 4 1000 | 0.491882 | 0.114127 | 4.31 |
| 5 3000 | 35.257888 | 2.655063 | 13.28 |

Implemented by C
using OpenMP and
blocking

ARGSORT

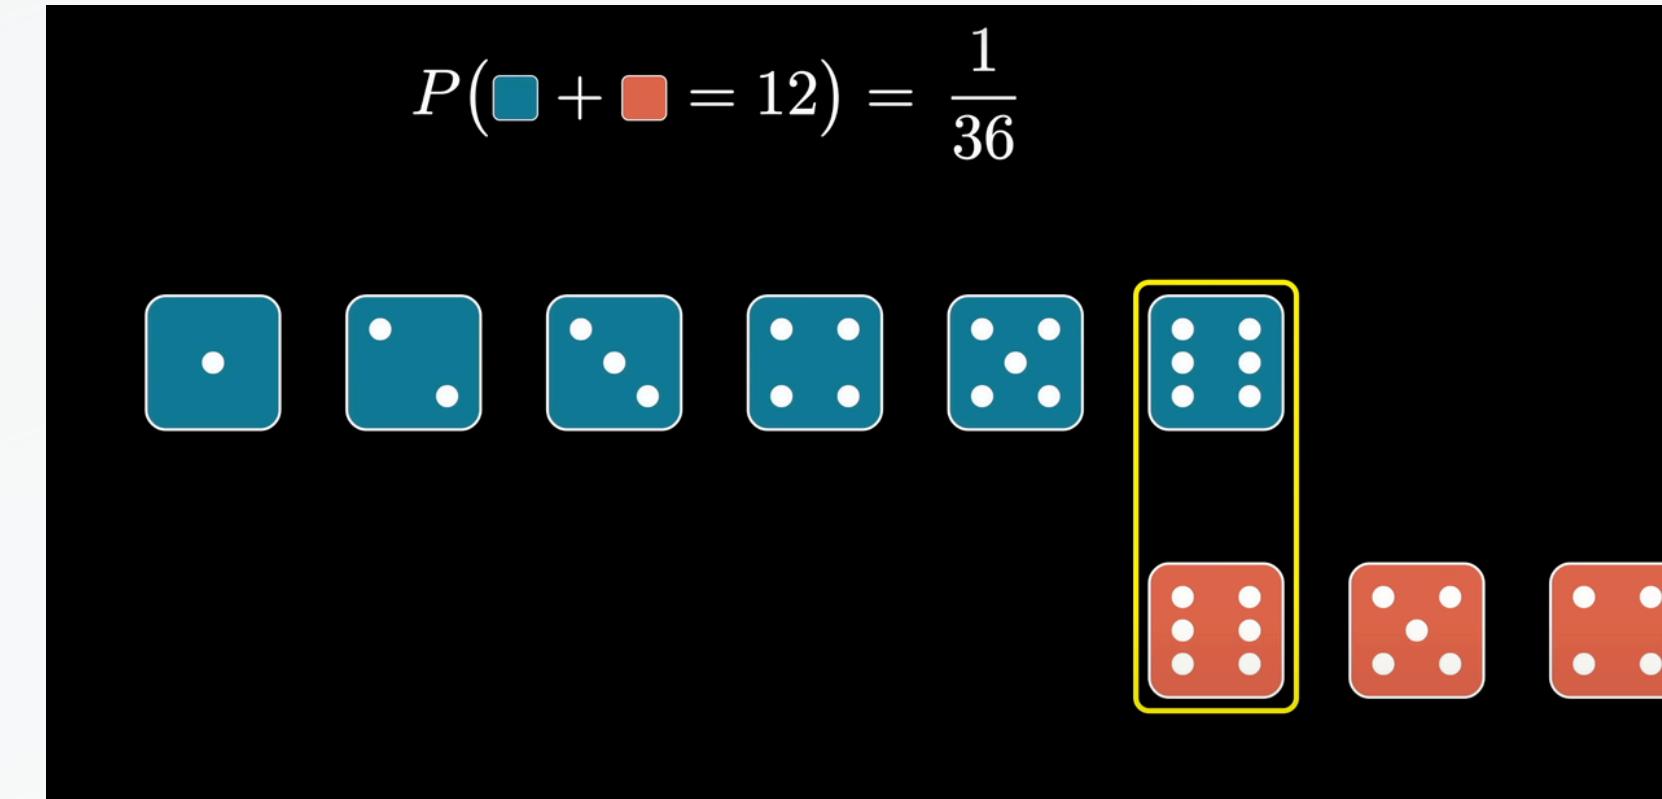
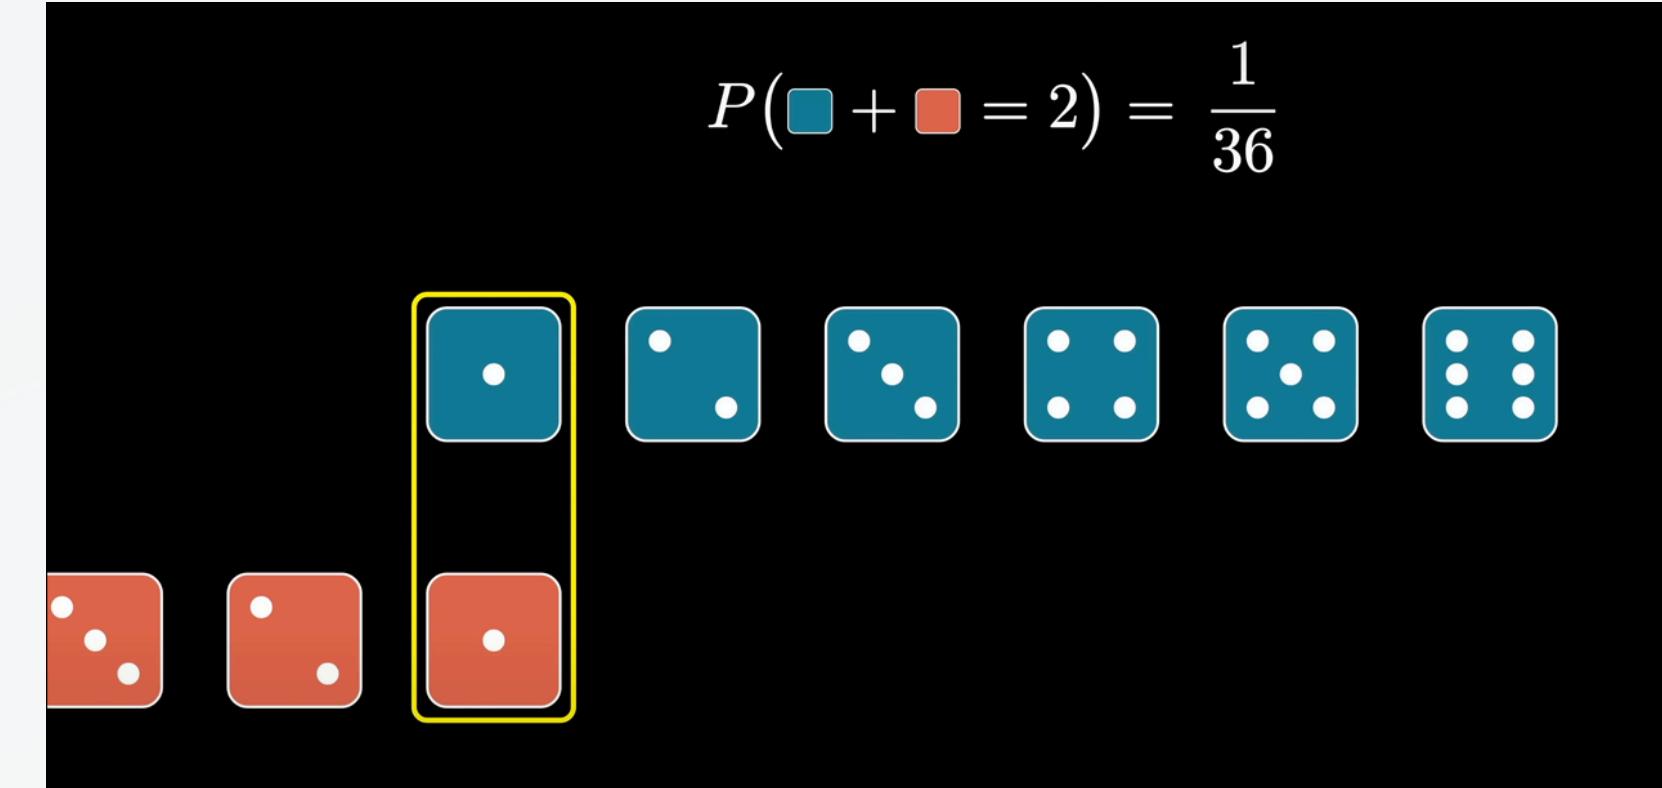
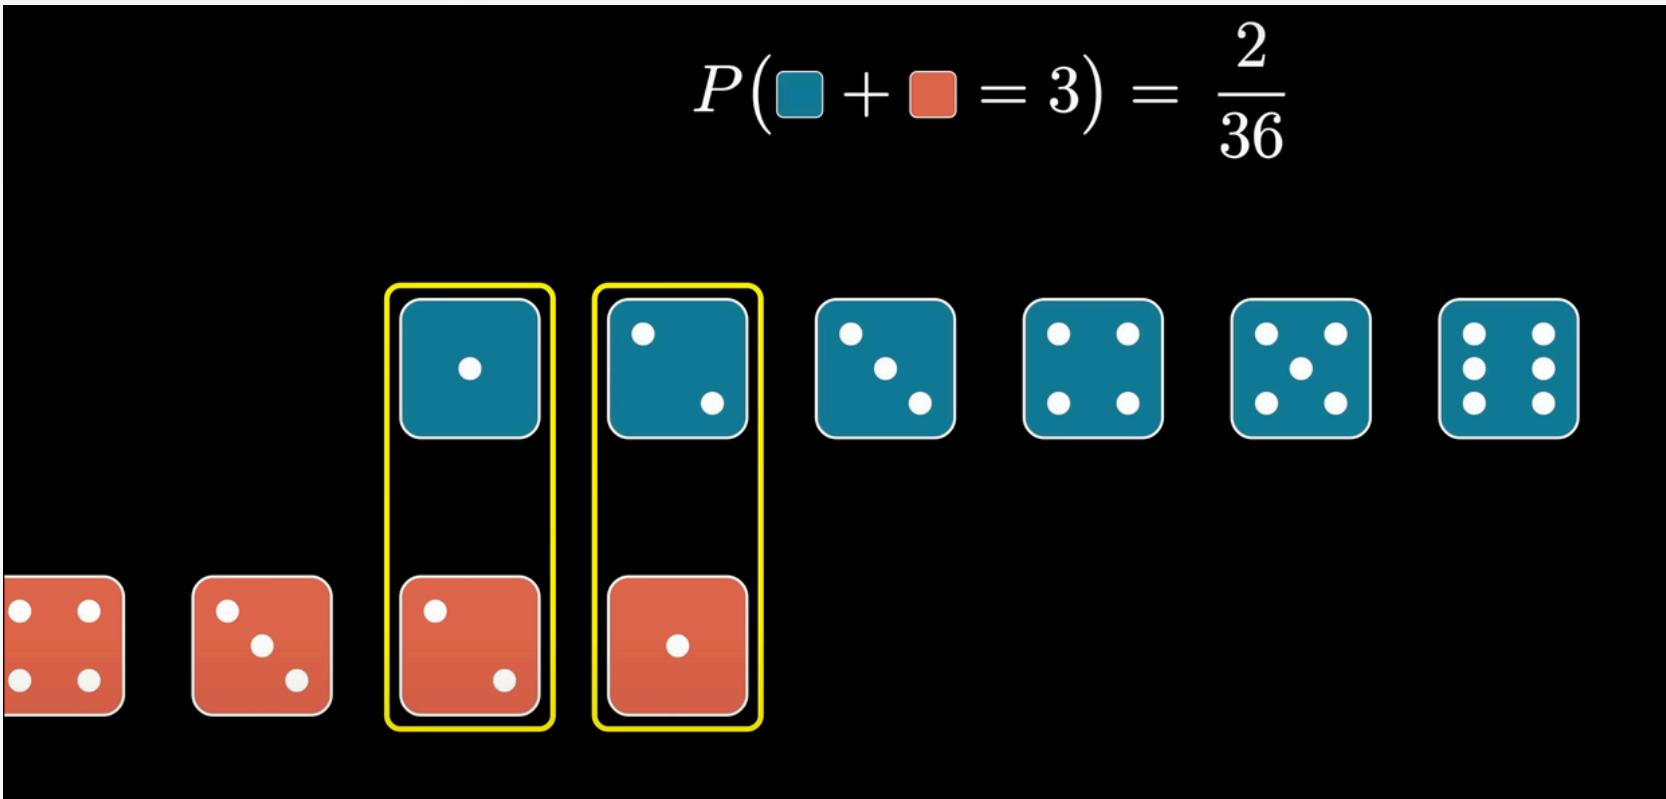
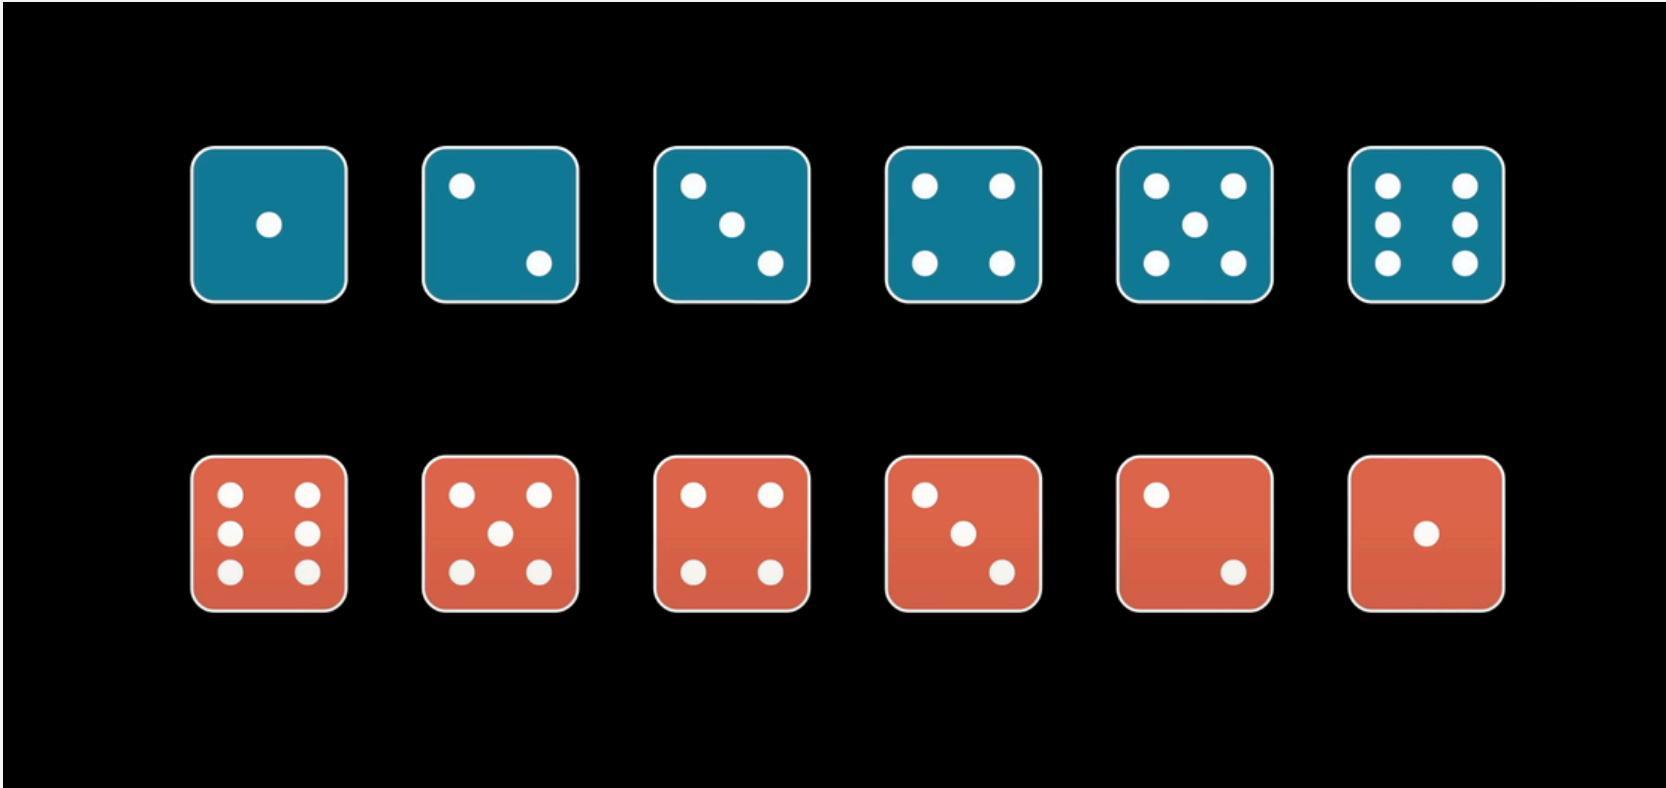


Implemented by
Rust using Rayon



| | Size | Numpy Time (s) | Rust Time (s) | Speedup Factor |
|---|-----------|----------------|---------------|----------------|
| 1 | 800 | 5.6e-05 | 5.2e-05 | 1.08 |
| 2 | 2000 | 0.000136 | 0.000139 | 0.98 |
| 3 | 10000 | 0.000709 | 0.000525 | 1.35 |
| 4 | 100000 | 0.00675 | 0.003369 | 2.0 |
| 5 | 5 million | 0.37189 | 0.192737 | 1.93 |

CONVOLVE



CONVOLVE

$$P(\text{blue} + \text{red} = 2) = a_1 \cdot b_1$$

$$P(\text{blue} + \text{red} = 3) = a_1 \cdot b_2 + a_2 \cdot b_1$$

$$P(\text{blue} + \text{red} = 4) = a_1 \cdot b_3 + a_2 \cdot b_2 + a_3 \cdot b_1$$

$$P(\text{blue} + \text{red} = 5) = a_1 \cdot b_4 + a_2 \cdot b_3 + a_3 \cdot b_2 + a_4 \cdot b_1$$

$$P(\text{blue} + \text{red} = 6) = a_1 \cdot b_5 + a_2 \cdot b_4 + a_3 \cdot b_3 + a_4 \cdot b_2 + a_5 \cdot b_1$$

$$P(\text{blue} + \text{red} = 7) = a_1 \cdot b_6 + a_2 \cdot b_5 + a_3 \cdot b_4 + a_4 \cdot b_3 + a_5 \cdot b_2 + a_6 \cdot b_1$$

$$P(\text{blue} + \text{red} = 8) = a_2 \cdot b_6 + a_3 \cdot b_5 + a_4 \cdot b_4 + a_5 \cdot b_3 + a_6 \cdot b_2$$

$$P(\text{blue} + \text{red} = 9) = a_3 \cdot b_6 + a_4 \cdot b_5 + a_5 \cdot b_4 + a_6 \cdot b_3$$

$$P(\text{blue} + \text{red} = 10) = a_4 \cdot b_6 + a_5 \cdot b_5 + a_6 \cdot b_4$$

$$P(\text{blue} + \text{red} = 11) = a_5 \cdot b_6 + a_6 \cdot b_5$$

$$P(\text{blue} + \text{red} = 12) = a_6 \cdot b_6$$

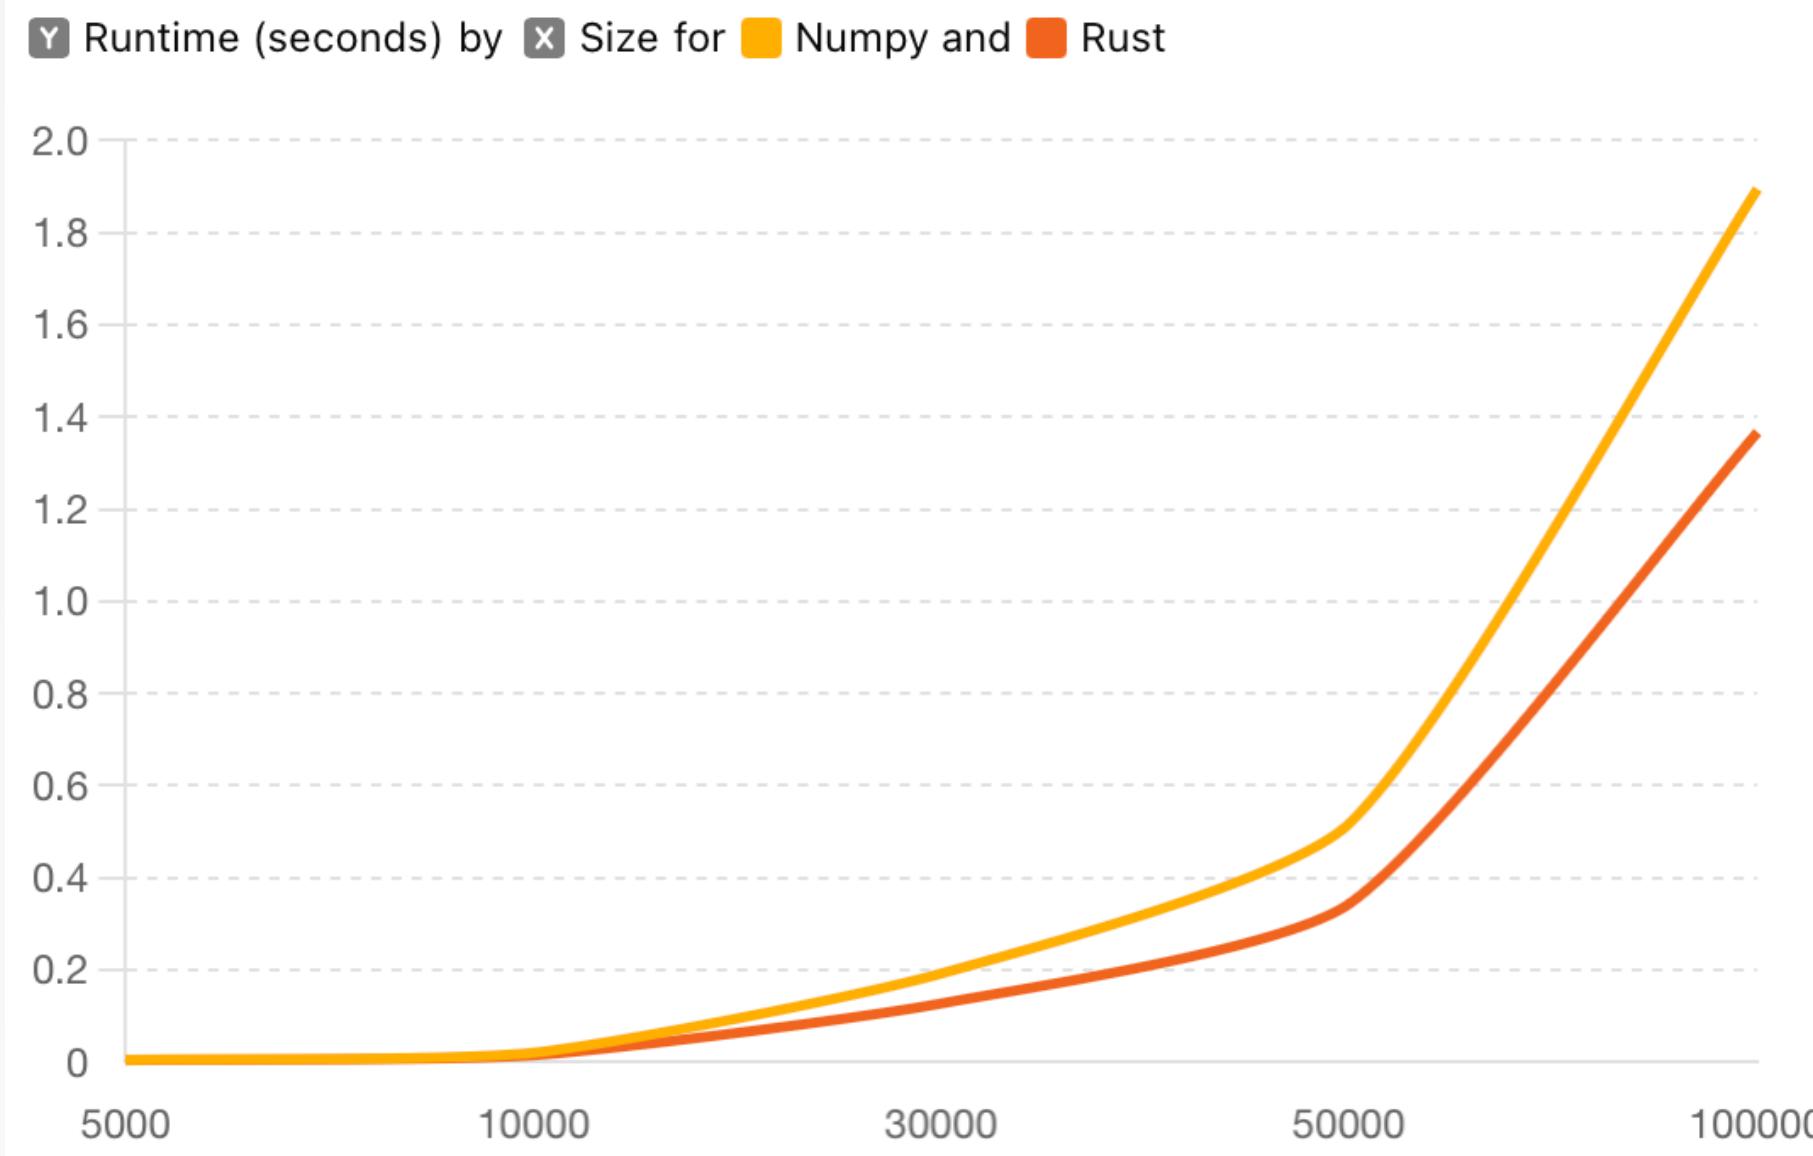
Convolution of (a_i) and (b_i)

$$(a * b)_n = \sum_{\substack{i,j \\ i+j=n}} a_i \cdot b_j$$



CONVOLVE

| | Size | Numpy Time (s) | Rust Time (s) | Speedup Factor |
|---|--------|----------------|---------------|----------------|
| 1 | 5000 | 0.004614 | 0.003895 | 1.18 |
| 2 | 10000 | 0.020123 | 0.014725 | 1.37 |
| 3 | 30000 | 0.192098 | 0.127037 | 1.51 |
| 4 | 50000 | 0.517631 | 0.343414 | 1.51 |
| 5 | 100000 | 1.89345 | 1.365702 | 1.39 |



Implemented by
Rust using Rayon
and thread pool

All simple taken
from M2 8-core cpu